



IS
2nd
material

Database 2




❖ TRANSACTION PROCESSING

SINGLE-USER vs MULTI-USER SYSTEMS:


• Single user system:

- It is designed to be used by only **one user** at a time.
- It is typically installed on a single computer and can only be accessed by the user who installed it or the user who is currently logged in.



• Multi user system:


- It can be accessed by **multiple** users simultaneously.
- It is typically installed on a network server and can be accessed by users who are logged in to the network.



INTERLEAVED vs PARALLEL PROCESSING:

- **Interleaved processing:** It is a process in which more than one task is being processed at **the same time**.

- **Parallel processing:** It is a process in which the tasks are divided into small sub-tasks that are processing **simultaneously or parallel**



A transaction :

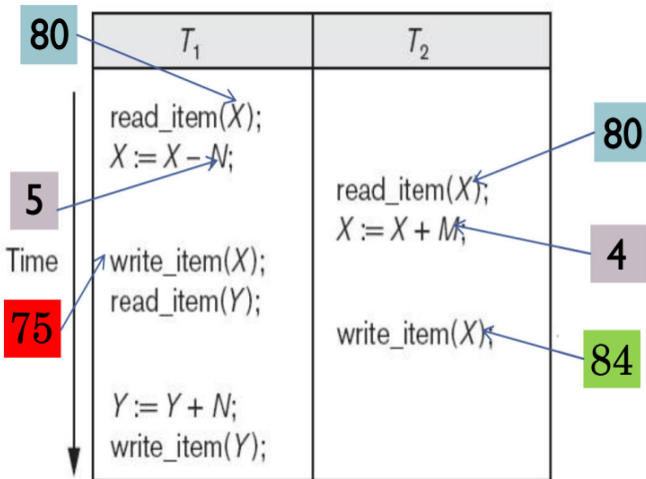
- it is the **logical unit** of database processing that includes one or more database access operations (Read: retrieval or Write: insertion, deletion, modification).
- It is the **atomic unit** of work that is either completed in its entirety or not done at all.
- In the application program it may contain several transactions separated by **transaction boundaries**.

IMPORTANCE OF CONCURRENCY CONTROL

- The lost update.
- The temporary update (**Dirty read**).
- The incorrect summary problem.
- The unrepeatable read problem.


Lost update problem: This problem occurs when **two transactions** that access the same database items have their operations **interleaved** in a way that makes the value of some database items incorrect.

اول مشكله عندنا هي ال **lost update** و دي بيقولك لما يكون ف اتنين **transactions** و الاتنين بيستخدمو نفس الداتا بيز بحيث ان عملياتهم **interleaved** يعني بيحصلو مع بعض ف نفس الوقت و مترتبين ع بعض و ركز ع الى علمتلك عليه ده عشان هنتاجه تحت ف المثال ف بيقولك ده ممكن يخلي ان قيم بعض العناصر الى ف الداتا بيز تطلع غلط .



مثال هيفهمك الدنيا شوية .. ف المثال الى ع اليمين ده عندك داتا بيز فيها اتنين **transactions** الى هم **T1** و **T1, T2** هيشغل الاول هيقرا قيمة المتغير **X** و بعدين يعمل عليها عملية و قبل م يسجل قيمه ال **X** بعد م اتغيرت راح قرا قيمتها القديمه و عمل عليها عملية تانيه ف ده نتج عنه انه لما خزن البيانات طلع قيمتين للاكس مختلفين ع بعض

Temporary update (Dirty Read): This problem occurs when **one transaction updates a database item** and then the **transaction fails** for some reason "Meanwhile, the updated item is accessed (read) by another transaction before it is changed back to its original value."



تاني مشكله عندنا هي ال **temporary update** و دي بتسمى بال **Dirty Read** بيقولك ان دي بتحصل لما **transaction** يعمل ابديت لعنصر و بعدين ال **transaction** ده يفشل او يكراش لسبب معين "الى زى م قولتلك فوق هنتكلم عنهم بعدين" المهم ان ف نفس الوقت العنصر الى اتعمله ابديت ده **transaction** تاني استخدمه بعدها علطول قبل م يرجع لقيمتة الاساسيه . مثال هيفهمك شويه الدنيا .. ف المثال الى ع اليمين ده عندك داتا بيز فيها اتنين **transactions** الى هم **T1** و **T1, T2** هيشغل الاول هيقرا قيمة المتغير **X** و بعدين يعمل عليها عملية و يسجله علطول (و ده الى محصلش ف اول مشكله) و بعدين

transaction تاني هيسخدم قيمه ال **X** الجديده الى اتسجلت خلاص و يعمل عليها عملية تانيه و يسجلها و بعدين حصل ف ال **transaction** الاول **Failure** و هو بيحاول يقرأ ال **X** ف ده هيقف ال **transaction** ده كله بس ف نفس الوقت زى م قولنا خلاص ف **transaction** تاني استخدم العنصر الى اتعمله ابديت

Incorrect summary problem: This problem occurs if **one transaction is calculating an aggregate summary function** on a number of database items **while other transactions are updating some of these items**, the aggregate function may calculate some values before they are updated and others after they are updated

T_1	T_3
<pre> read_item(X); X := X - N; write_item(X); read_item(Y); Y := Y + N; write_item(Y); </pre>	<pre> sum := 0; read_item(A); sum := sum + A; : : read_item(X); sum := sum + X; read_item(Y); sum := sum + Y; </pre>

تالت مشكله و دى ابسط نوع فيهم هى ال *incorrect summary* ف المشكله دى بكل بساطه بتحصل لما يكون عندك اتنين *transactions* واحد منهم فيه داله ال *sum* بتحسب مجموع عنصر معين و ف نفس الوقت العنصر ده بيتعمل لقيمه ابديت ف *transaction* تانى ف ده بيخليك يبقى ف قيمتين لداله ال *sum* واحده قبل الابديت و واحده بعد الابديت . مثال .. ف المثال الى ع اليمين ده عندك داتا بيز فيها اتنين *transactions* الى هم T_1, T_2 و T_2 هيعمل عمليه ال *sum* للمتغير X الاول قبل م يحصل ابديت لاي عنصر بعدها علطول ف ال *transaction* الاول هيجعل ابديت لقيمه ال X ف لما هيرجع بعدها لل *transaction* التانى عشان يعمل عمليه *sum* تانيه هيبقى عنده قيمه جديده مختلفه لل X بعد الابديت ف ده الى بيقولك عليه هيبقى ف قيمتين مختلفين لل *sum* نتيجته ان ال X هيبقى ليها هى كمان قيمتين مختلفين واحده بعد الابديت و التانيه قبل

Unrepeatable read problem: This problem occurs when a transaction T **reads the same item twice** and **the item is changed by another transaction T between the two reads**. Hence, T receives different values for its two reads of the same item.

اخر مشكله عندك بيكلّمك ع ال *unrepeatable read problem* و الى هى بتحصل لما يكون ف نفس ال *transaction* بيقرا قيمه ال X مرتين بس بين المرتين دول ف *transaction* تانى عمل ابديت لقيمه ال X و متنساش احنا قولنا ان كل دول متوصلين مع بعض *interleaved* و انت فاهم بقى خلاص يعنى ايه .. ف لما هيجعل ابديت بين العمليتين هيخلي ف نفس ال *transaction* يبقى ف قيمتين لل X مختلفين ع

IMPORTANCE OF RECOVERY

- It means that the system is responsible for making sure that:
 - All the operations in the transaction are **completed successfully** and their **effect** is recorded permanently in the database.
 - The transaction **does not have any effect** on the **database** or **any other transactions**. (means that there is no transaction with failure that will effect the database or any other transaction)

REASONS FOR TRANSACTION FAILURE:

- **Computer failure:** If the hardware crashes, the contents of the computer's internal memory may be lost

• بيحصل لما ال **hardware** يكراش او مثلا ان محتوى الذاكرة الداخليه اتمسح

- **Transaction or system error:** Transaction operation error (integer overflow or division by zero). Erroneous parameter values or a logical programming error. The **user** may interrupt the transaction during its execution

• بيحصل لما مشكله ف العمليات الى بتتعمل جوا ال **transaction** زي مثلا انه يكتب رقم غلط او يقسم ع صفر

او لما تدخل قيمه متغير غلط او يبقى ف **logical programming error**

او لما اليوزر نفسه هو الى يوقف ال **transaction** و هو بيتنفذ

- **Local errors:** Data for the transaction may not be found. A programmed abort in the **transaction** causes it to fail

• بيحصل لما مثلا الداتا بتاعت ال **transaction** الجهاز او البرنامج مش عارف يلاقيها

او لما يحصل **abort** جوا ال **transaction** نفسه ف يخليه يحصله **fail**.

- **Concurrency control enforcement:** The **concurrency control method** may decide to abort the transaction, to be **restarted later**.

• بيحصل لما ال **concurrency control method** هي الى توقف ال **transaction** عشان تعمله ريستارت

- **Disk failure:** disk read/write head crash.
- **Physical problems and catastrophes:** Ex. power or air-conditioning failure, fire, theft, sabotage, or overwriting disks.

TRANSACTION AND SYSTEM CONCEPTS

- A transaction is the **atomic unit** of **work** that is either completed in its entirety or not done at all.
- For **recovery** purposes, the system needs to keep track of when the transaction **starts**, **terminates**, and **commits** or **aborts**.

TRANSACTION STATES:

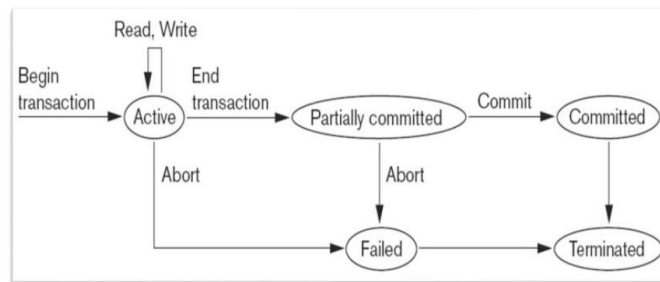
- Active state
- Terminated state.
- Partially committed state.
- Committed state.
- Failed state.

Recovery manager keeps track of the following operations:

- **Begin-transaction:** This marks the **beginning** of transaction execution.
- **Read or Write:** These specify **read** or **write** operations on the database items that are executed as part of a transaction.

End-transaction: This specifies that **read** and **write** transaction operations **have ended** and **marks the end limit** of transaction execution.

- At this point it may be necessary to check whether **the changes introduced by the transaction** can be permanently **applied to the database** or whether **the transaction** has to be **aborted** because it violates concurrency control or for some other reason.



ببقولك في ال **end-transaction** بيحدد امتى ال **read and write transaction** خلصت
و بيحدد كمان اخر transaction execution

- و كمان ف المرحلة دى بيبقى لازم تحدد اذا كان التغيرات الى عملتها ال **transaction** اتنفذت بشكل صحيح و هتقدر تطبق ع الداتا بيز او ان ال **transaction** لازم يتعمله **abort** عشان فيه اى مشكله م الى اتكلمنا عنهم ف ال **concurrency control** او لاي سبب تانى

Commit-transaction: This signals a **successful end** of the transaction so that any changes (updates) executed by the transaction can be safely committed to the database and **will not** be undone.

- **Rollback (Abort):** This signals that the transaction **has ended unsuccessfully**, so that any changes or effects that the transaction may have applied to the database **must be** undone.

SYSTEM LOG:

- It keeps track of **all transaction operations** that affect **the values of database items**.
- It may be needed to permit **recovery** from **transaction failures**.
- The log is kept on disk, so it **is not affected** by any type of failure **except** for **disk or catastrophic failure**.
- The log is periodically **backed up to archival storage (tape)** to guard against such **catastrophic failures**.

Commit point of a transaction:

- A transaction T reaches its commit point when **all its operations** that access the database have been **executed successfully** and **the effect** of all the transaction operations on the database has been **recorded in the log**.
- After **the commit point**, the transaction **is said to be committed**, and **its effect** is assumed to be **permanently recorded** in the database.

بيقولك ان ال transaction يوصل لل commit point لما كل العمليات الخاصه بال transaction تنتفذ بنجاح و ان التأثير بتاع كل العمليات ع الداتا بيز يتحفظ ف ال system log
و كمان بعد م ال transaction يكون وصل لل commit point بيتقال عليه انه committed و ان تأثيره او التغييرات الى حصلت ف الداتا بيز نتيجته العمليات الى عملها اتحفظت بشكل دائم ف الداتا بيز

ACID properties:

Atomicity: A transaction is an **atomic unit of processing**; it is either performed in its entirety or not performed at all.

Consistency preservation: A correct execution of the transaction must take the database **from one consistent state to another**.

Isolation: A transaction **should not make its updates visible to other transactions** until it is committed; this property when enforced strictly, **solves the temporary update problem**, and makes **cascading rollbacks of transactions unnecessary**

Durability (permanency): Once a transaction changes the database and the changes are committed, these changes must never be lost because of **subsequent failure**.

Transaction schedule or history: It is the order of execution of operations from the various transactions, when transactions are executing concurrently in an **interleaved** fashion.

- It is an **ordering of the operations of the transactions** subject to the constraint that, for each transaction T_i that participates in S , the operations of T_i in S must appear **in the same order** in which they occur in T_i .

ايه هو ال Transaction schedule

اول تعريف عندك بيقولك انه هو عمليه ترتيب العمليات م ال transactions المختلفه حسب تنفيذها و ده لما ال transactions تكون بتنفذ عملياتها بطريقه ال interleaved الى اتكلمنا عليها قبل كده .

تاني تعريف بيقولك انها ترتيب العمليات بتاعت ال transactions حسب ال constraint او القواعد بتاعت كل transaction بمعنى ان كل transaction T_i بيشارك ف ال schedule S العمليات بتاعته لازم تكون مكتوبه بنفس الترتيب الى بيحصل ف ال transaction

SCHEDULES CLASSIFIED ON RECOVERABILITY:

- **Recoverable schedule:** In it **no** transaction needs to **be rolled back**.
- A schedule S is recoverable if **no transaction** T in S **commits** until all transactions T' that have written an item that T reads have **committed**.
- **Cascadeless schedule:** In it **every transaction reads only** the items that are written by committed transactions.
- **Schedules requiring cascaded rollback:** A schedule in which **uncommitted transactions** that read an item from a failed transaction **must be rolled back**.

- **Strict schedule:** A schedule in which a transaction can **neither** read nor write an item **X** until the last transaction that wrote **X** has committed.
- **Serial schedule:** A schedule **S** is serial if, for every transaction **T** participating in the schedule, **all the operations of T** are executed **consecutively** in the schedule.
- **Serializable schedule:** A schedule **S** is serializable if it is **equivalent** to some **serial schedule** of the same **n** transactions.

Result equivalent: Two schedules are called result equivalent **if they produce the same final state** of the database.

Conflict equivalent: Two schedules are said to be conflict equivalent **if the order of any two conflicting operations is the same in both schedules**.

Conflict serializable: A schedule **S** is said to be conflict serializable **if it is conflict equivalent to some serial schedule S'**.

- Being **serializable** means that the schedule is a correct schedule which means that: It will leave the database in **a consistent state**.

Interleaving is appropriate and will result in a state as if the transactions were **serially executed**, and it will achieve efficiency due to **concurrent execution**.

- Serializability is hard to check due to:
 - Interleaving of operations **occurs in an operating system** through some scheduler.
 - It is difficult to determine beforehand **how the operations** in a schedule **will be interleaved**.
- It's **not possible** to determine when a schedule begins and when it ends.

Conflict occurs if two operations satisfy with:

- They belong to **different** transactions.
- They access the **same** item **X**.
- At least one of the operation is **write_item(X)**.
- EX: • $r_2(X)$ and $w_1(X)$: **conflict**
 - $r_1(X)$ and $r_2(X)$: **Not conflict**
 - $w_2(X)$ and $w_1(Y)$: **Not conflict**
 - $r_1(X)$ and $w_1(X)$: **Not conflict**

View equivalence schedule: A **less restrictive definition** of equivalence of schedules.

View serializability schedule:

- Definition of serializability **based on view equivalence**.
- A schedule is view serializability **if it is view equivalent to a serial schedule**.

Two schedules are said to be **equivalent** if the following three conditions hold:

- **The same set of transactions participates in S and S'**, and **S and S' include the same operations of those transactions**.

For any operation $R_i(X)$ of T_i in **S**, if the value of **X** read by the operation has been written by an operation $W_j(X)$ of T_j (or if it is the original value of **X** before the schedule started), the same condition must hold for the value of **X** read by operation $R_i(X)$ of T_i in **S'**.

- If the operation $W_k(Y)$ of T_k is **the last operation to write Y in S**, then $W_k(Y)$ of T_k must also be **the last operation to write item Y in S'**.

بيقولك امتى يبقى **Two schedules** متكافئين؟

اول حاجه لما يبقى نفس ال **Transactions** موجوده ف الاتنين يعنى لو عندى **Transaction** بيعمل مثلا عمليه جمع لمتغير معين ف ال **S** يبقى لازم يبقى ف **Transaction** بيعمل عمليه جمع ف ال **S'** .. و بيقولك كمان ع انه مش يبقى فيهم نفس ال **Transaction** و خلاص لا زى م قولتك قبلها لو واحد بيعمل جمع يبقى التانى يعمل جمع بردو.

طيب ركز معايا ف تانى حاجه لو عندى مثلا **Transaction** ف ال **S** بيعمل **Read** للمتغير **X** بيقولك ان لو قيمه المتغير **X** الى اتعملها **Read** ف اول **Transaction** اتعملها **Write** ف **Transaction** تانى ف ال **S** بردو .. ف كل الى حصل ده هيطبق زى م هو كده ع قيمه المتغير **X** الى اتعملها **Read** ف ال **Transaction** الى ف ال **S'**

تالت حاجه بيقولك لو عندنا **Transaction** بيعمل **Write** للمتغير **Y** و عمليه ال **Write** دى هى اخر عمليه ف ال **S** يبقى لازم بردو تبقى هى اخر عمليه ف ال **S'**.

The premise behind view equivalence:

- As long as **each read operation** of a transaction reads the result **of the same write operation** in both schedules, the write operations of each transaction must produce the same results.
- "The view"**: the read operations are said to see the same view in both schedules.

Relationship between view and conflict equivalence:

- The two are same under constrained write assumption** which assumes that if T writes X, it is constrained by the value of X it read.
- Conflict serializability is stricter than view serializability. With unconstrained write (or blind write), a schedule that is view serializable is not necessarily conflict serializable.**
- Any **conflict** serializable schedule **is also view** serializable, but not vice versa.

Consider the following schedule of three transactions:

- T₁: R₁(X), W₁(X); T₂: W₂(X); T₃: W₃(X);
- S_a: R₁(X); W₂(X); W₁(X); W₃(X); C₁; C₂; C₃;
- Since T₂ and T₃ do not read the value of X, then the operations W₂(X) and W₃(X) are **blind writes**.
- S_a is **view serializable**, since it is view equivalent to the serial schedule T₁, T₂, T₃.
- However, **S_a is not conflict serializable** since it is not conflict equivalent to any serial schedule.
- **Under special semantic constraints**, schedules that are otherwise **not conflict serializable** may work **correctly**.
- **Using commutative operations of addition and subtraction** (which can be done in any order) certain **non-serializable** transactions may work **correctly**.